

Working with daily climate model output data in R and the **futureheatwaves** package

by G. Brooke Anderson, Colin Eason, Elizabeth A. Barnes

Abstract Research on climate change impacts can require extensive processing to incorporate output from ensembles of different climate models and different simulations of each model. This processing can be particularly extensive when identifying and characterizing extreme events like heat waves and frost day spells, as these must be processed from model output with daily time steps and may use identification thresholds based on local distributions (e.g., 98th percentile temperature for that location). Here, we provide an overview of working with daily climate model output data in R. Further, we present the **futureheatwaves** package, which we developed to ease the process of identifying, characterizing, and exploring multi-day extreme events in climate model output. This package takes as input a directory of projection files from one or many climate models, identifies all occurrences of a specific type of extreme event (e.g., heat waves) for multiple cities using a user-specified event definition, and writes an output directory of files with these extreme events and their characteristics (e.g., length, absolute and relative temperatures). The **futureheatwaves** package then allows the user to create custom functions and apply them across all event files; this functionality can be used either to explore extreme event patterns (e.g., average length and intensity across different climate models) or to apply epidemiological estimates to project health impacts under different scenarios.

Introduction

Research on climate change impacts can require extensive processing. This is not only because output files for a single climate model can be large, but also because of the rising popularity of ensemble techniques (Cubasch et al., 2013), in which, to better characterize uncertainty in projections, impacts are assessed for multiple climate models, multiple simulations of each climate model, and multiple climate experiments. Projections of regional climate change over the next century are subject to high uncertainty due to three distinct sources: (1) internal climate variability, i.e. climate noise, (2) climate model uncertainty, i.e. the same forcing can produce a different response in different models and (3) scenario uncertainty, i.e. uncertainty in future climate forcings (e.g. Hawkins and Sutton (2009)). One approach for ensuring that these sources of uncertainty are characterized is to simulate the future climate many times with multiple models and for multiple future scenarios.

For example, the Coupled Model Intercomparison Project, phase 5 (CMIP5; Taylor et al. (2012)) brought together dozens of major climate modeling groups around the world to simulate the same future radiative forcing scenarios, but with their own models. This created an ensemble of state-of-the-art climate model projections that allow researchers to study projections and their uncertainties. Most of these modeling groups additionally performed more than one simulation for each scenario and model (i.e. multiple ensemble members), perturbing the initial conditions by a very tiny amount to quantify uncertainties due to internal climate variability.

Processing climate output data is particularly intensive for data in smaller time steps, like daily output. While some climate impacts can be assessed using climate model output at monthly, seasonal, or yearly time steps, the impacts of multi-day extreme events must be assessed using output in daily time steps. Such multi-day extreme events include heat waves, cold spells, frost day spells, and droughts. The assessment is further complicated if extreme events are identified based on conditions that are rare for a certain location (e.g., 98th percentile of local temperature distribution for identifying heat waves) (Cubasch et al., 2013), as this event definition must be determined at each study location from climate model output. Further, it is often of interest to create summaries of multiple characteristics of these extreme events. For example, one may be interested in determining whether the frequency or characteristics (e.g., length, intensity) of heat waves or warm spells will change under certain climate change scenarios (Cubasch et al., 2013).

We begin this article with an overview of climate model output data, particularly daily data, for R users. We outline where data from CMIP5 can be obtained as well as how to work with the file format (netCDF) from R. We overview some R packages that can be useful when working with this data, as well as aspects of the data (e.g., non-standard calendars) of which users should be aware when working with daily climate model output in R.

After this overview, we present the **futureheatwaves** package which we created to aid in identifying and characterizing any type of multi-day extreme event from daily climate model output

Design goals of the futureheatwaves package	
1	Make processing of large sets of climate projections more practical for researchers exploring the potential impacts of heat waves and other multi-day extreme events.
2	Speed up processing time by incorporating C++ in event identification.
3	Keep track of the names of climate models, and number of ensemble members processed for each.
4	Not only identify, but also characterize, all extreme events within each climate projection, to allow the exploration of patterns in these characteristics across different projections and also to allow the use of more complex impacts models, including models that include effect modification by event characteristics (e.g., event length, event intensity). For example, this package allows the user to apply a health effects model where risk of mortality is not the same for every heat wave, but rather is modified by heat wave length, intensity, or other measured characteristics.
5	Give users extensive power in customizing the process, including allowing custom event definitions.
6	Allow users to easily explore the extreme events identified within all climate projections by applying custom functions across heat wave data sets from all projections at once.
7	Create output that is in a “tidy” data format, allowing it to work well with ggplot2 for visualization.

Table 1: Design goals for the **futureheatwaves** package.

(Table 1). This package can be used to identify and characterize a variety of multi-day extreme events, but provides some functionality particularly useful in identifying and characterizing heat waves specifically. Quantification of the impacts of heat waves on human health suffer from additional sources of uncertainty beyond those inherent in projections of regional changes in surface temperature. These include: (1) uncertainty in the ability of communities to adapt to changing temperatures (the adaptation scenario) and (2) uncertainty in the definition of a heat wave itself. Thus, identifying and characterizing the impacts of future heat waves in state-of-the-art climate models requires analyzing hundreds of projections based on combinations of anthropogenic activity, climate model and ensemble member, adaptation scenario, and heat wave definition. Such an analysis is non-trivial. For example, to estimate the possible impacts of heat waves using 32 ensemble members of CMIP5, two forcing scenarios, two different heat wave definitions, and five progressive assumptions about adaptation to heat, would require that one identify and characterize heat waves in 640 separate temperature time series for each region of interest.

The **futureheatwaves** package automates the process of identifying and characterizing multi-day extreme events across different ensemble members of one or more climate models. A variety of different heat wave definitions have been used to identify heat waves in a time series of temperature data (Smith et al., 2013), and the choice of heat wave definitions can influence both health effect estimates (Chen et al., 2015; Kent et al., 2014) and projected heat wave trends (Smith et al., 2013). Further, other types of extreme events will be defined differently than heat waves (for example, frost day spells may be defined as one or more days with temperature at or below 32°F / 0°C). This package therefore allows the user to create and use a custom extreme event definition to identify events in the climate model output. Finally, this package allows the user to explore the extreme events identified in the climate model output with a function that can take a user-defined R function and apply it across all extreme event files generated for the separate ensemble members.

An overview of climate model output for R users

CMIP5 climate model output data

For climate impact studies, a main source for climate model output is the Coupled Model Intercomparison Project, which is currently in its fifth phase (CMIP5). Over 20 climate modeling groups have created one or more climate models which, for this project, are run using standardized scenarios

(Taylor et al., 2012). The resulting output is uniform across modeling groups and has a consistent structure, which allows comparison of simulations from different models (Flato et al., 2013). For CMIP5, each group ran simulations under several experiments, with experiments varying in terms of the radiative forcing. This radiative forcing depends on time-varying model inputs (greenhouse gas emissions or concentrations, land use changes, etc.), which are specified for each experiment (Taylor et al., 2012; Flato et al., 2013). Experiments include historical experiments (run using forcings consistent with observed and reconstructed data for 1850–2005), pre-Industrial control experiments, and experiments of future scenarios of radiative forcing over the 21st century or longer (e.g., RCP4.5, RCP8.5) (Taylor et al., 2012). Some modeling groups created ensembles of output for a specific model and experiment, in which they ran the experiment multiple times with the model with very small changes to the initial conditions, resulting in multiple ensemble members for a single climate model and experiment. CMIP5 climate model output is archived at a number of different time steps (e.g., daily, monthly, seasonal, yearly) (Taylor and Doutriaux, 2010), and some variables are reported at multiple levels in the ocean or atmosphere (e.g., ocean temperature). Here, we will focus on data with a daily time step for variables reported at a single level (e.g., near-surface air temperature).

The CMIP5 climate model output data is distributed across data nodes at different climate modeling centers (Taylor et al., 2012), but can be accessed centrally at the World Climate Research Programme CMIP5 data portal at <https://pcmdi.llnl.gov/search/cmip5/>. Users must register before downloading data, and some data are restricted to non-commercial use. There is a separate file for each combination of climate model, experiment, modeling realm (e.g., atmosphere, ocean), variable, time step, and ensemble member (Taylor et al., 2012; Taylor and Doutriaux, 2010). For finer time scales, this output is further split across multiple files for specific year ranges (e.g., 5 years of output for each file) (Taylor and Doutriaux, 2010). Filenames for CMIP5 files can be parsed to generate information about the output variable, climate model, experiment, and ensemble member for the simulation (Taylor and Doutriaux, 2010). Files can be searched and downloaded through a point-and-click web interface. The data portal also allows you to generate a specific `wget` script, which you can use to download many files at once. A `wget` script can also be generated through Earth System Grid Federation's Search RESTful API. Tips on efficiently searching and downloading the data, including through use of `wget` scripts and the search API, are available as user tutorials through the website of the University of Colorado Boulder's Earth System CoG (e.g., <https://www.earthsystemcog.org/projects/cog/doc/wget> for a tutorial on downloading files using `wget`).

CMIP5 files are saved in Network Common Data Format (netCDF). NetCDF is a binary file format that allows storage of data representing a regular array. These climate model output files can be as large as several gigabytes [?]. For climate model output at a single level (e.g., near-surface air temperature), the data is a 3-dimensional array, with dimensions representing time and two coordinates of location (e.g., latitude and longitude) (Figure 1). Global climate models generate output at regularly-spaced time steps, typically at regularly-spaced grid points around the world. The latitude and longitude spacing of grid points vary by climate model, but are typically 1–2 degrees for atmospheric variables in CMIP5 models (Flato et al., 2013). Each data point in the netCDF array gives the modeled value of the variable (e.g., surface temperature) for a single time point and location. For CMIP5 climate model output, the location units are in degrees east and degrees north for longitude and latitude, respectively. For daily output files, the time unit is in days since a specified origin date-time (e.g., days since 1850-01-01 00:00:00) (Taylor and Doutriaux, 2010). All CMIP5 output files are required to include certain metadata (Taylor and Doutriaux, 2010). This required metadata includes the experiment, forcing agents input to the model to create the simulation, time step, institution and institutional contact information, climate model, and modeling realm (Taylor and Doutriaux, 2010). The metadata also must include units for all of the coordinate variables (e.g., longitude, latitude, time).

Taylor et al. (2012) and Meehl et al. (2007) are excellent resources for finding out more about the CMIP climate model output data.

Working with climate model output in R

A few R packages can be used to work with the netCDF file format used for CMIP5 files. Earlier packages to work with netCDF files included `ncdf` and `ncvar`, but these do not work with the newer netCDF version 4 released in 2008 and are no longer available through CRAN. More recent packages, including `ncdf4` (Pierce, 2015) and `RNetCDF` (Michna and Woods, 2013, 2016), work with both version 4 and netCDF's older version 3. Climate model output data for CMIP5 is required to conform with the earlier version (version 3) (Taylor and Doutriaux, 2010) and so should work with any of these packages, although it is safer to write code using functions that can be used with either version in case future phases of CMIP do not require files to conform with netCDF version 3.

The netCDF format allows you to access metadata and variables describing the dimensions of the data without reading the full file into memory. The metadata describes the dimensions of the

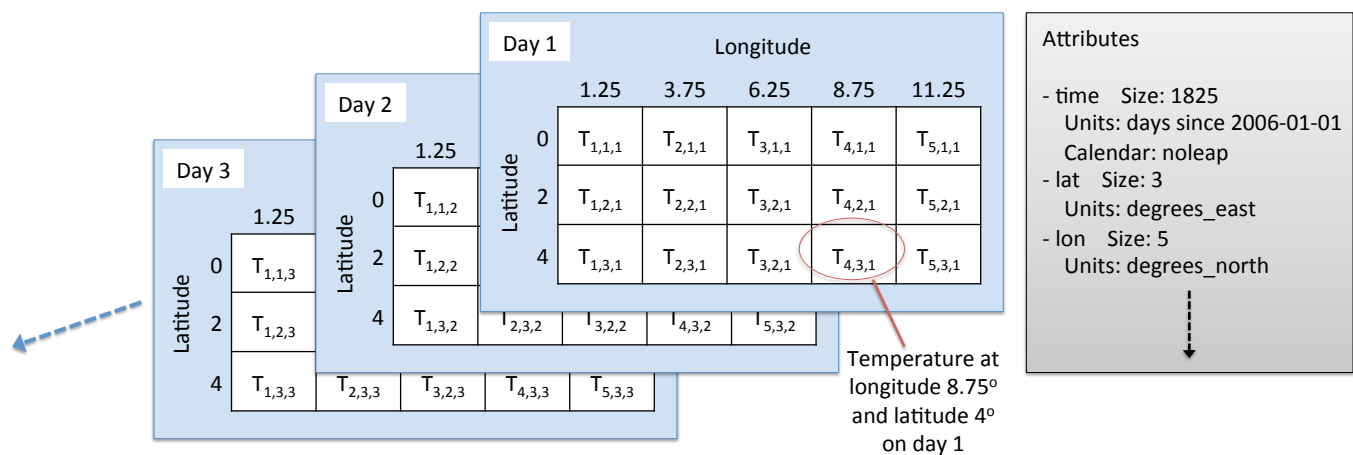


Figure 1: Example of structure of a NetCDF climate model output file for a variable reported at a single depth, like surface air temperature. Data are stored in a three-dimensional array, with measurements at each time step and grid location. Surface temperature data are typically indexed in climate model output files by longitude, latitude, and time, in that order. For example, if the air surface temperature ("tas") is read into an R object called `tas`, you can access the value for the first day at the fourth longitude and third latitude with `tas[4, 3, 1]`. In addition to the output variable (temperature in this example), vectors with the ordered values of each dimension (longitude, latitude, and time) can also be read in from the netCDF file, as well as attribute data (e.g., units for variables, the calendar used for time).

array and the units, etc., of each variable and is printed by the print method of the R object returned by the `nc_open` function of **ncdf4**. For CMIP5 files, the dimension variables include the times and locations corresponding to each array in the data. Before reading in a variable from a netCDF file, a connection to the file must be opened, for example with the `nc_open` function from **ncdf4**. Variables can then be read in using the `ncvar_get` function from **ncdf4**, with the `varid` parameter set to "lat", "lon", or "time" (as a caveat, many climate models output to non-Gregorian calendars, in which case the time variable should be read in using a different function, as discussed later in this section). The climate output variable (e.g., near-surface air temperature) can similarly be read in using `ncvar_get`. In this case, the `varid` parameter should be set using the appropriate CMIP5 variable name (e.g., "tas" for near-surface air temperature); these variable names can be found in the CMIP Requested Output tables (Taylor and Doutriaux, 2010). You can use the dimensional time and location data to identify the location of the variable data you need in the netCDF array and use indexing to read only that data into memory, without needing to read in the full file (for example, with the `nc.get.var.subset.by.axes` function in **ncdf4.helpers**). Once the user is done reading in data from the file, the connection can be closed (e.g., with the `nc_close` function from **ncdf4**).

Since the late 1500s, Western dates have been set using the Gregorian calendar, which has 365.2425-day years. Some climate models, however, are run using different calendars, including the Julian calendar (365.25-day years), a calendar where there are no leap years (365-day years), a calendar where every year is a leap year (366-day years), and a calendar of twelve 30-day months (360-day years) (Eaton et al., 2011). With these non-Gregorian calendars, R's base functions for converting a vector to a Date class based on the number of days since an origin date (`as.Date`, `as.POSIXct`) do not return the desired values. The **PCICt** (Bronaugh and Drepper, 2013) and **ncdf4.helpers** (Bronaugh, 2014) packages provide further functionality with netCDF files that can be useful when working with climate model output data and provide particular help in working with different calendars. CMIP5 netCDF files include information on the calendar used, which the `nc.get.time.series` function in **ncdf4.helpers** use to convert the "time" variable in the file to an object of the **PCICt** class, which provides Date-like functionality for 360- and 365-day calendars (Bronaugh and Drepper, 2013). As a note, while these functions will help with handling most CMIP5 files, the CMIP5 standards allows use of other calendars which may not be successfully handled by these functions, so it is important to assess whether the time variable range in the **PCICt** object correctly matches the expected date ranges for a file as you process CMIP5 data in R.

The size of the climate model output files can be large enough that it may make more sense to work with smaller chunks of the data, rather than reading all data into memory and working with the data all at once (Todd-Brown and Bond-Lamberty, 2016). This problem aggregates when working with multiple climate models and more than one ensemble member for each of those climate models.

Modeled temperature on a day in July 2075

GFDL-ESM2G model, RCP8.5 experiment, r1i1p1 ensemble member

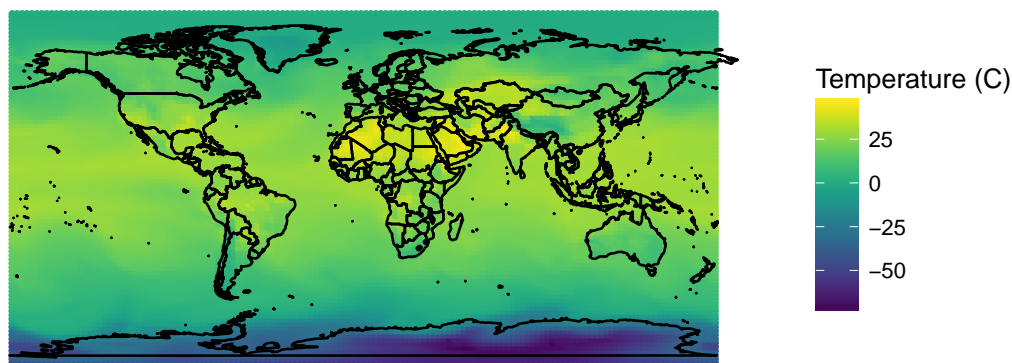


Figure 2: Example of mapping near-surface air temperature data worldwide for a single day of climate model output data. This map uses data from the Geophysical Fluid Dynamics Laboratory’s Earth System Model 2G, r1i1p1 ensemble member, on a single day in the summer of 2075.




Figure 3: Example of plotting a time series of temperature simulations between 2071 and 2075 from CMIP5 daily climate model output data for the model grid cell point closest to Beijing, China. This plot uses data from the Geophysical Fluid Dynamics Laboratory’s Earth System Model 2G, r1i1p1 ensemble member.

The functions and packages described can be used to work with CMIP5 netCDF files in R to do things like map near-surface air temperatures from a single climate model on a specific day in July 2075 (Figure 2) and pull a time series of daily near-surface air temperature simulations for the climate model grid point closest to Beijing, China (Figure 3). The “starting_from_netcdf” vignette that comes with the “futureheatwaves” package provides all code required to create these figures, as well as more details and code examples on working with CMIP5 netCDF files in R.

In addition to these general packages for working with netCDF files, there are several R packages specifically for working with climate model output data, including **RCMIP5** (Todd-Brown and Bond-Lamberty, 2016) and **wux** (Mendlik et al., 2016). However, these packages are more useful for working with data output at time steps of a month or higher and have limited utility with the daily climate model output data required for studies of multi-day extreme events.

The **RCMIP5** package includes functions to read in CMIP5 data from netCDF files, scan a directory of CMIP5 files and determine models with continuous available data, create objects of a special “cmip5data” class to work with CMIP5 data within R, and parse the filenames for all files in a directory to extract information within the filename. For this package, most functions only work with monthly or less frequent data (e.g., the functions `checkTimePeriod` and `cmip5data`, as well as functions that work with `cmip5data` objects including `filterDimensions` and `getProjectionMatrix`) (Todd-Brown




Figure 4: Overview of the functionality of the **futureheatwaves** package. The package takes a directory with climate projection files (left), for one or more climate models, with one or more ensemble members for each climate model (this example figure shows four climate models with one or two ensemble members each). The `gen_hw_set` function processes these files to create a data frame for each ensemble member, identifying and characterizing all multi-day extreme events (e.g., heat waves) in the time series projection for that ensemble member. The `apply_all_models` function allows users to explore these extreme events by applying user-created functions across all the extreme event data frames, creating a summary data frame with results.

and Bond-Lamberty, 2016). While the `loadCMIP5` function does successfully load the daily data as a `cmip5data` object, most of the methods for this object type do not do anything meaningful for daily data. The package’s `getFileInfo` function, however, will work with CMIP5 files of any time step; this function identifies all CMIP5 files in a directory and creates a dataframe that parses the information contained in the file name. As a note, the `get.split.filename.cmip5` function in the **ncdf4.helper** package similarly can be used to parse information contained in CMIP5 file names (Bronaugh, 2014).

The **wux** package (Mendlik et al., 2016) includes functions that allow the user to download CMIP5 monthly-aggregated output directly from within R with the `CMIP5fromESGF` function. However, this function does not allow downloading of climate model output with finer time steps, like daily data. This package uses the `models2wux` function to read in climate model output netCDF files and convert to a WUX dataframe that can be used by other functions in the package. While this function can input climate model output with daily time steps, if the element “`what.timesteps`” of the `modelinput` list input is set to “`daily`”, the function aggregates this data to a monthly or less frequent (e.g., seasonal) aggregation when creating the WUX dataframe. Therefore, while this package provides very useful functionality for working with averaged output of daily climate model output data, it cannot easily be used to identify and characterize multi-day extreme events like heat waves.

More tips on working with CMIP5 files in R are provided in the “`starting_from_netcdf`” vignette of the **futureheatwaves** package.

The futureheatwaves package

How the package works

We created the **futureheatwaves** package to aid in identifying, characterizing, and exploring multi-day extreme events in daily climate model output data. Figure 4 gives an overview of the two primary functions of the **futureheatwaves** package. First, the `gen_hw_set` function processes a directory of climate projection files that are stored locally on the user’s computer (Figure 4, “Climate projections”), to generate a list of all extreme events in each projection, as well as over a dozen characteristics of each identified extreme event (Table 2). The extreme events are identified and characterized at one or more study locations (e.g., cities), with study locations specified by the user in an input file. The extreme events identified for each ensemble member are output as separate files in a directory specified by the user (Figure 4, “Extreme events datasets”).

Once the user creates these data frames of location-specific heat waves, the `apply_all_models` function allows the user to apply custom functions across all the extreme event data frames generated by `gen_hw_set` to summaries of extreme events across all climate models and ensemble members (Figure 4, right). This functionality can be used to generate summary statistics (e.g., determine average heat wave length or total frost days), or can be used to apply more complex functions (e.g., apply epidemiologic effect estimates across the heat waves to generate health impact estimates).

Example data

We have included data files in the package to serve as example files so that users can try this package before applying it to their own directory of climate projection files. This example data comes from two climate models that are a part of CMIP5: (1) the model of the Beijing Climate Center, China

Column name	Description of characteristic
mean.var	Average daily value of the variable across all days in the extreme event, in the units in which the variable is expressed in input files (e.g., average daily mean temperature during the heat wave in degrees Kelvin)
max.var	Highest daily value of the variable across all days in the extreme event, in the units in which the variable is expressed in input files
min.var	Lowest daily value of the variable across all days in the extreme event, in the units in which the variable is expressed in input files
length	Number of days in the event
start.date	Date of the first day of the event
end.date	Date of the last day of the event
start.doy	Day of the year of the first day of the event (1 = Jan. 1, etc.)
start.month	Month in which the event started (1 = January)
days.above.abs.thresh.1	Number of days in the event above a specified absolute threshold (default is the number of days in the event above 80°F / 26.7°C, but this can be changed with the absolute_thresholds argument in gen_hw_set)
days.above.abs.thresh.2	Number of days in the event above a specified absolute threshold (default is the number of days in the event above 85°F / 29.4°C)
days.above.abs.thresh.3	Number of days in the event above a specified absolute threshold (default is the number of days in the event above 90°F / 32.3°C)
days.above.abs.thresh.4	Number of days in the event above a specified absolute threshold (default is the number of days in the event above 95°F / 35.0°C)
days.above.99th	Number of days in the event above the 99 th percentile of the variable for the location, using the period specified with the referenceBoundaries argument in gen_hw_set as a reference for determining these percentiles
days.above.99.5th	Number of days in the event above the 99.5 th percentile of the variable for the location, using the period specified with the referenceBoundaries argument in gen_hw_set as a reference for determining these percentiles
first.in.year	Whether the event was the first to occur in its calendar year in the location
mean.var.quantile	The percentile of the average variable value during the event compared to the location's year-round distribution of the variable, based on the variable distribution for the location during the period specified by the referenceBoundaries argument in gen_hw_set
max.var.quantile	The percentile of the maximum variable value during the event compared to the location's year-round distribution of the variable, based on the variable distribution for the location during the period specified by the referenceBoundaries argument in gen_hw_set
min.var.quantile	The percentile of the minimum variable value during the event compared to the location's year-round distribution of the variable, based on the variable distribution for the location during the period specified by the referenceBoundaries argument in gen_hw_set
mean.seasonal.var	The location's average seasonal value of the variable (by default, season is set to May–September, but this can be changed with the seasonal_months argument in gen_hw_set), based on the variable values for the location during the years specified by the referenceBoundaries argument in gen_hw_set
mean.yearround.var	The location's average year-round value of the variable, based on the variable values for the location during the years specified by the referenceBoundaries argument in gen_hw_set

Table 2: Extreme event characteristics measured by the `gen_hw_set` function in the **futureheatwaves** package. The left column gives the name of each variable's column in the extreme event datasets created by the `gen_hw_set` function. When characterizing extreme events below a threshold, like cold spells, appropriate alternatives are given for some columns (e.g., `days.below.abs.thresh.1`, `days.below.1st`).

Meteorological Administration (BCC) (Xin et al., 2013) and (2) the National Center for Atmospheric Research's (NCAR's) Community Climate System Model, version 4 (CCSM4) (Gent et al., 2011). We include one ensemble member from BCC (r1i1p1) and two from CCSM (r1i1p1 and r2i1p1).

To ensure that the size of this example data is reasonably small, we have only included projection data for grid points from these climate models that are near five U.S. east coast cities: New York, NY; Philadelphia, PA; Newark, NJ; Baltimore, MD, and Providence, RI. Further, to keep the file sizes reasonably small, the historical projections range over the years 1990 to 1999, while the future projections are limited to 2060 to 2079. Users' applications of this package will likely use directories with many more climate model ensemble members and more locations; however, the operation of the package is the same for this smaller example application as it would be for a much larger application.

Once the **futureheatwaves** package is installed and loaded, the user can find the local location of these files using R's `system.file` function. In the later sections of this article, we show how to use the package functions with these example files as inputs.

When using this package with data other than this example data, the climate model output data will require some pre-processing. Data will need to be saved in a specific format, with files stored in a specific directory structure. This preprocessing can be done using some of the R packages described earlier for working with CMIP5 data in R. Full details of the required file and directory structure are provided in the package's "futureheatwaves" vignette, while tips on conducting this processing starting from CMIP5 netCDF files are given in the "starting_from_netcdf" vignette.

Basic example of identifying and characterizing extreme events

Once climate model output files are set up as specified in the "futureheatwaves" package vignette, the package can process them to identify and characterize heat waves in each ensemble member's projection for each location using the `gen_hw_set` function. For example, to process the example climate model output data included with the package, the user can run:

```
projection_dir_location <- system.file("extdata/cmip5",
                                     package = "futureheatwaves")
city_file_location <- system.file("extdata/cities.csv",
                                 package = "futureheatwaves")

gen_hw_set(out = "example_results",
           dataFolder = projection_dir_location,
           dataDirectories = list("historical" = c(1990, 1999),
                                "rcp85" = c(2060, 2079)),
           citycsv = city_file_location,
           coordinateFileNames = "latitude_longitude_NorthAmerica_12mo.csv",
           tasFileNames = "tas_NorthAmerica_12mo.csv",
           timeFileNames = "time_NorthAmerica_12mo.csv")
```

This code first identifies and saves as objects the path names on the user's computer of the example climate projections directory (`projection_dir_location`) and the file of study locations (`city_file_location`). The `gen_hw_set` function processes this example input and creates a new directory, 'example_results', with files of identified and characterized heat waves, in the user's current working directory. In this example code, the processing is done using default values for the event definition, years for which to generate the heat wave data sets, etc. How and why to customize these choices are explained later in the text. Function arguments (e.g., `dataDirectories`, `tasFileNames`) are used to specify the format of the data and the directory structure.

Once the function has completed running, results will be written locally to the directory specified by the `out` argument of `gen_hw_set`. This directory will include files with some basic information about the climate models and the closest grid points of each climate model to each location, as well as a directory with files of identified and classified extreme events for each ensemble member, including all characteristics in Table 2.

Customizing the extreme event definition

By default, the package identifies extreme events in climate model output data using a specific definition for heat waves:

A *heat wave* is two or more days at or above a city-specific threshold temperature, with the threshold determined as the 98th percentile of year-round temperature in the city during some reference period (by default, 1990–1999).

However, the user can extensively customize the definition used to identify extreme events. Users can write a custom R function with either a different heat wave definition (see [Smith et al. \(2013\)](#) and [Kent et al. \(2014\)](#) for listings of definitions used in scientific studies) or with a definition appropriate for a different type of extreme event (e.g., one or more days at or below 32°F / 0°C for frost day spells). For heat wave identification, researchers might want to use a specific definition, for example, because it matches the definition used by local health officials to declare heat wave warnings or, in the case of health impact assessments, to match with a definition used in an epidemiological study. For other extreme events, a heat wave definition likely will not be applicable and so a customized definition is necessary.

Three components of the extreme event definition can be easily customized in the `gen_hw_set` function call, without creating a new R function to use to identify heat waves. First, often users wish to define an extreme event based on conditions that are rare in the study location ([Cubasch et al., 2013](#)). In this case, the percentile of the variable of interest required for an extreme event can be changed using the `probThreshold` option in `gen_hw_set`. For example, this option could be used to identify heat waves based on the 99th percentile, rather than the 98th percentile, of a city's year-round temperature. Second, the user can change the number of days used in the heat wave definition using the `numDays` argument in the `gen_hw_set` function (the default is `'numDays = 2'`, corresponding to two or more days). Combined, these two customization choices allow the user to identify heat waves using many of the heat wave definitions used in previous climate and health research— for example, 9 of the 16 heat wave definitions used in [Kent et al. \(2014\)](#) (e.g., [some examples of the 9 that could be fit]) could be fit using different combinations of these two options for specifying threshold percentile and number of days. Third, some extreme events like cold waves and frost day spells are defined as a certain number of days below, rather than above, a threshold. While the default is to identify events by searching for days above a threshold, this behavior can be changed with the `above_threshold = FALSE` argument in the `gen_hw_set` function.

Beyond these simpler options, the customization of the event definition is even more extensive as one has the option of writing a custom R function. This functionality allows the user to use heat wave definitions that either require a number of days above an absolute threshold (e.g., maximum temperature of $\geq 95^{\circ}\text{F}$ for ≥ 1 day [Kent et al. \(2014\)](#); [Tan et al. \(2007\)](#); minimum temperature $\leq 0^{\circ}\text{C}$ for ≥ 1 day for frost day spells) or that require a combination of thresholds to be met (e.g., maximum daily temperature one threshold every day of the heat wave and above a higher threshold for a certain number of days; [Kent et al. \(2014\)](#); [Peng et al. \(2011\)](#)). To use a customized event definition, the user can write and load an R function that implements the definition. The custom function is passed to the `gen_hw_set` function using the `IDheatwavesFunction` argument. To work correctly, this custom function must allow only specific inputs and generate only specific outputs. Details about this required structure are provided in the **futureheatwaves** package vignette. To increase processing speed when identifying extreme events, we coded parts of the default event definition function in C++ and synced it with R using the **Rcpp** package ([Eddelbuettel and Francois, 2011](#)). Users should consider a similar strategy for custom heat wave definitions, especially if processing a large number of climate projection files.

Finally, extreme events are often based on conditions that are rare for a specific location using location-specific relative thresholds. These thresholds are often defined for climate impact studies based on a variable's distribution at that location in present-day or historical data. However, it can be interesting to explore trends in extreme events under climate change if extreme events are identified based on variable distributions during the projection period, or another future period. It is possible to specify the time period to use when determining a relative threshold for an event definition using the `thresholdBoundaries` argument in the function `gen_hw_set`.

Exploring extreme events

Once the user has created a directory of characterized event files for each ensemble member ("Extreme event data sets", Figure 4), he or she can explore the results using the `apply_all_models` function. This function allows the user to apply custom R functions across all extreme event data frames created by the `gen_hw_sets` call. The user can apply any R function that follows certain standards in accepting input and returning output. Full details on these standards are given in the **futureheatwaves** package vignette.

As an example, if the user wanted to calculate the average temperature of the heat waves identified for each ensemble member, he or she could write a simple function:

```
average_mean_temp <- function(hw_datafr){
  out <- mean(hw_datafr$mean.var)
  return(out)
}
```

This function could then be applied across all extreme event data sets output by `gen_hw_set` using the `apply_all_models` function. For example, to apply this function to all the example output results that come with the package, the user could run:

```
out <- system.file("extdata/example_results", package = "futureheatwaves")
apply_all_models(out = out, FUN = average_mean_temp)

#>   model ensemble   value
#> 1  bcc1         1 302.3745
#> 2  ccsm         1 302.4458
#> 3  ccsm         2 302.3428
```

This output gives the results (value column) of running the custom function for each ensemble member of each climate model. Note that the location of the directory with the heat wave data frames must be specified using the `out` argument when calling `apply_all_models`. Typically, this will be the directory path for the directory specified with the `out` argument in `gen_hw_set`.

Location-specific results can be generated using the `city_specific` argument in `apply_all_models`:

```
apply_all_models(out = out, FUN = average_mean_temp, city_specific = TRUE)

#>   model ensemble city   value
#> 1  bcc1         1 balt 305.1816
#> 2  bcc1         1  nwk 300.3367
#> 3  bcc1         1   ny 300.3367
#> 4  bcc1         1 phil 305.1816
#> 5  bcc1         1 prov 298.0402
#> 6  ccsm         1 balt 303.1277
#> 7  ccsm         1  nwk 302.4053
#> 8  ccsm         1   ny 302.4053
#> 9  ccsm         1 phil 302.3425
#> 10 ccsm         1 prov 301.8895
#> 11 ccsm         2 balt 302.9373
#> 12 ccsm         2  nwk 302.2748
#> 13 ccsm         2   ny 302.2748
#> 14 ccsm         2 phil 302.2858
#> 15 ccsm         2 prov 301.9520
```

This output is structured as “tidy” data (Wickham, 2014), allowing it to be used easily with the graphing package `ggplot2` (Wickham, 2009).

The `apply_all_models` function can also be used to project the health impacts of heat waves. A simple estimate of excess deaths associated with heat wave risk in a community can be calculated as (Peng et al., 2011):

$$E_c = B_c * (RR - 1) * L_c \quad (1)$$

where:

- E_c is the total number of excess deaths in community c ;
- B_c is the baseline average daily mortality in community c ;
- RR is the relative risk of mortality associated with a heat wave day; and
- L_c is the total number of heat wave days in community c over the study period.

This impact assessment equation can be translated into a function that merges each projection’s heat wave data frame with a data frame of community-specific baseline mortality rates (B_c), calculates equation 1 for each heat wave, and then sums up the total excess deaths across all heat waves. In this example, we have set the default relative risk for a heat wave as 1.032, based on the estimated added heat wave relative risk in 107 U.S. communities for heat waves defined as two or more days at or above a community’s 98th percentile temperature (Anderson and Bell, 2009).

```
excess_deaths <- function(hw_datafr, base_mortality, RR = 1.032){
  hw_datafr <- dplyr::left_join(hw_datafr, base_mortality,
                                by = "city") %>%
    dplyr::mutate(excess_deaths = base_mort * (RR - 1) * length)
  out <- sum(hw_datafr$excess_deaths)
```

```

    out <- round(out)
    return(out)
}

```

Once defined in R, this function can be applied across all heat waves from all climate models' ensemble members, provided that you have a data frame called `city_mortality` with columns with each community's identifier (`city`) and baseline mortality rate (`base_mort`), using the call (this example uses hypothetical baseline mortality rates for the cities and requires the [dplyr](#) package (Wickham and Francois, 2016)):

```

library(dplyr)
city_mortality <- data.frame(city = c("balt", "nwk", "ny", "phil", "prov"),
                             base_mort = c(3, 2.5, 35, 6, 2.5))
apply_all_models(out = out, FUN = excess_deaths, base_mortality = city_mortality)

#>   model ensemble value
#> 1  bcc1         1   930
#> 2  ccsm         1   968
#> 3  ccsm         2   967

```

While the above impact analysis is simplistic, it provides an example of how the `apply_all_models` function can be used in conjunction with epidemiologic results to estimate the potential health impacts of extreme events with climate change.

Customizing dates for projections and reference temperatures

Depending on the climate model output data used and the research question, different projection periods may be of interest. For example, one researcher may wish to explore extreme events in the 2050–2075 period, while another may be interested in 2090–2100. The projection date range can be changed by specifying starting and ending years in the `projectionBoundaries` argument of the `gen_hw_set` function.

Further, some of the event characteristics (e.g., `mean.temp.quantile`, Table 2) are calculated based on relative temperature, which is a measure of how the value of the variable of interest during an extreme event compares to the typical distribution of that variable at that location. These characteristics are measured for each of the extreme events identified by the `gen_hw_set` function by taking the absolute value of the variable during the event (e.g., average temperature during the heat wave is 90°F 32.2°C) and comparing it to the location's typical variable distribution. This process generates relative measures of how intense the event is compared to what is normal in that location (e.g., 90°F 32.2°C is in the 99th percentile of year-round temperatures in the location).

These relative event characteristics will vary depending on whether you calculate them based on a location's present-day variable distribution or on the location's variable distribution in the future, since the distributions of many relevant variables (e.g., temperature, precipitation) are expected to change in many locations with climate change. The package therefore allows the user to specify date ranges of the temperature distributions to be used in calculating these relative temperature metrics in each location, which can be done using the `referenceBoundaries` option of `gen_hw_set`.

Mapping grid points

It can be useful to explore the location of the climate model grid point used to pull climate model output for each study location with a given climate model. Therefore, the package has a function called `map_grid_leaflet` that plots the locations of grid points used for each location from each climate model. This function is built using the `htmlWidget` [leaflet](#) package (Cheng and Xie, 2016). The following code illustrates the use of this function with the example data to create Figure 5, which plots the grid points used in the example data from the BCC climate model in the example data:

```

out <- system.file("extdata/example_results", package = "futureheatwaves")
map_grid_leaflet(plot_model = "bcc1", out = out)

```

Extensions

While this package was created to be used for research on extreme events in climate change projections, it can be used more broadly. For example, there are other episodes like wildfires and air pollution where it may be interesting to identify extended periods of high exposures in projection time series, and this package could be applied to gridded air pollution model output to explore these exposures.

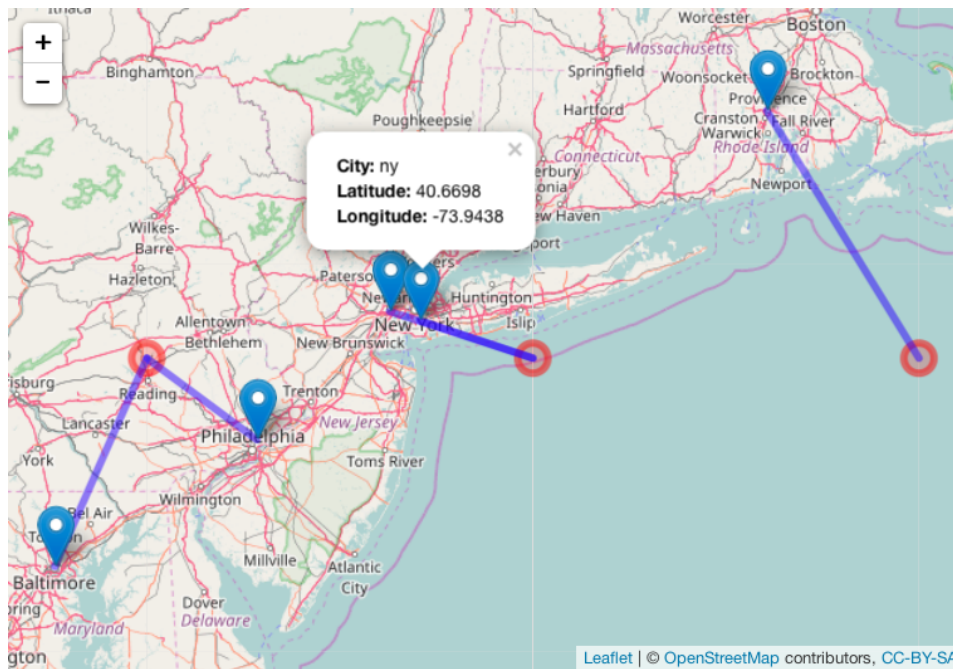


Figure 5: Snapshot of an example interactive map created using the `map_grid_leaflet` showing the locations of study cities and their matching climate model grid points for the BCC climate model example data included with **futureheatwaves**. The lines on the map connect each climate model grid point to the study location(s) for which that grid point was used. The interactive maps include pop-ups with city identifiers; one is shown open in this snapshot as an example. From this map, you can see that the climate model grid point closest to New York City for this climate model is over the Atlantic Ocean.

Acknowledgements

This work was supported by grants from the National Institute of Environmental Health Sciences (R00ES022631), the National Science Foundation (1331399), and the Colorado State University Vice President for Research. Thank you to early testers of the package: Julia Bromberek, Wande Benka-Coker, Josh Ferreri, Ryan Gan, Molly Gutilla, Mike Lyons, Casey Quinn, Rachel Severson, and Meilin Yan.

Bibliography

- G. B. Anderson and M. L. Bell. Weather-related mortality: how heat, cold, and heat waves affect mortality in the United States. *Epidemiology*, 20(2):205–213, 2009. doi: 10.1097/EDE.0b013e318190ee08. [p10]
- D. Brunaugh. *ncdf4.helpers: Helper functions for use with the ncdf4 package*, 2014. URL <https://CRAN.R-project.org/package=ncdf4.helpers>. R package version 0.3-3. [p4, 6]
- D. Brunaugh and U. Drepper. *PCICt: Implementation of POSIXct work-alike for 365 and 360 day calendars.*, 2013. URL <https://CRAN.R-project.org/package=PCICt>. R package version 0.5-4. [p4]
- K. Chen, J. Bi, J. Chen, X. Chen, L. Huang, and L. Zhou. Influence of heat wave definitions to the added effect of heat waves on daily mortality in Nanjing, China. *Science of the Total Environment*, 506-507:18–25, 2015. doi: 10.1016/j.scitotenv.2014.10.092. [p2]
- J. Cheng and Y. Xie. *leaflet: Create Interactive Web Maps with the JavaScript ‘Leaflet’ Library*, 2016. URL <https://CRAN.R-project.org/package=leaflet>. R package version 1.0.1. [p11]
- U. Cubasch, D. Wuebbles, D. Chen, M. Facchini, D. Frame, N. Mahowald, and J.-G. Winther. Introduction. in: Climate change 2013: The physical science basis. contribution of working group i to the fifth assessment report of the intergovernmental panel on climate change. *Climate Change 2013*, 5: 119–158, 2013. [p1, 9]

- B. Eaton, J. Gregory, B. Drach, K. Taylor, S. Hankin, J. Caron, R. Signell, P. Bentley, G. Rappa, H. Höck, et al. NetCDF Climate and forecast (CF) metadata conventions version 1.6. 2011. URL <http://cfconventions.org/Data/cf-conventions/cf-conventions-1.6/build/cf-conventions.html>. [p4]
- D. Eddelbuettel and R. Francois. Rcpp: Seamless R and C++ integration. *Journal of Statistical Software*, 40(8):1–18, 2011. doi: 10.18637/jss.v040.i08. [p9]
- G. Flato, J. Marotzke, B. Abiodun, P. Braconnot, S. C. Chou, W. J. Collins, P. Cox, F. Driouech, S. Emori, V. Eyring, et al. Evaluation of climate models. in: Climate change 2013: The physical science basis. contribution of working group i to the fifth assessment report of the intergovernmental panel on climate change. *Climate Change 2013*, 5:741–866, 2013. [p3]
- P. R. Gent, G. Danabasoglu, L. J. Donner, et al. The Community Climate System Model Version 4. *Journal of Climate*, 24(19):4973–4991, 2011. doi: 10.1175/2011JCLI4083.1. [p8]
- E. Hawkins and R. Sutton. The potential to narrow uncertainty in regional climate predictions. *Bulletin of the American Meteorological Society*, 90(8):1095–1107, 2009. doi: 10.1175/2009BAMS2607.1. [p1]
- S. T. Kent, L. A. McClure, B. F. Zaitchik, T. T. Smith, and J. M. Gohlke. Heat waves and health outcomes in Alabama (USA): the importance of heat wave definition. *Environmental Health Perspectives*, 122(2): 151–158, 2014. doi: 10.1289/ehp.1307262. [p2, 9]
- G. A. Meehl, C. Covey, T. Delworth, M. Latif, B. McAvaney, J. F. B. Mitchell, R. J. Stouffer, and K. E. Taylor. The WCRP CMIP3 Multimodel Dataset: A new era in climate change research. *Bulletin of the American Meteorological Society*, pages 1383–1394, 2007. [p3]
- T. Mendlik, G. Heinrich, and A. Leuprecht. *wux: Wegener Center Climate Uncertainty Explorer*, 2016. URL <https://CRAN.R-project.org/package=wux>. R package version 2.2-1. [p5, 6]
- P. Michna and M. Woods. Rnetcdf—a package for reading and writing netcdf datasets. *The R Journal*, 5: 29–36, 2013. [p3]
- P. Michna and M. Woods. *RNetCDF: Interface to NetCDF Datasets*, 2016. URL <https://CRAN.R-project.org/package=RNetCDF>. R package version 1.8-2. [p3]
- R. D. Peng, J. F. Bobb, C. Tebaldi, L. McDaniel, M. L. Bell, and F. Dominici. Toward a quantitative estimate of future heat wave mortality under global climate change. *Environmental Health Perspectives*, 119(5):701–706, 2011. doi: 10.1289/ehp.1002430. [p9, 10]
- D. Pierce. *ncdf4: Interface to Unidata netCDF (Version 4 or Earlier) Format Data Files*, 2015. URL <https://CRAN.R-project.org/package=ncdf4>. R package version 1.15. [p3]
- T. T. Smith, B. F. Zaitchik, and J. M. Gohlke. Heat waves in the United States: definitions, patterns and trends. *Climatic Change*, 118(3-4):811–825, 2013. doi: 10.1007/s10584-012-0659-2. [p2, 9]
- J. Tan, Y. Zheng, G. Song, L. S. Kalkstein, A. J. Kalkstein, and X. Tang. Heat wave impacts on mortality in Shanghai, 1998 and 2003. *International Journal of Biometeorology*, 51(3):193–200, 2007. doi: 10.1007/s00484-006-0058-3. [p9]
- K. E. Taylor and C. Doutriaux. Cmp5 model output requirements: File contents and format, data structure and metadata. 2010. [p3, 4]
- K. E. Taylor, R. J. Stouffer, and G. A. Meehl. An overview of CMIP5 and the experiment design. *Bulletin of the American Meteorological Society*, 93(4):485–498, 2012. doi: 10.1175/BAMS-D-11-00094.1. [p1, 3]
- K. Todd-Brown and B. Bond-Lamberty. *RCMIP5: Tools for Manipulating and Summarizing CMIP5 Data*, 2016. URL <https://cran.r-project.org/web/packages/RCMIP5/>. R package version 1.2.0. [p4, 5]
- H. Wickham. *ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York, 2009. ISBN 978-0-387-98140-6. URL <http://had.co.nz/ggplot2/book>. [p10]
- H. Wickham. Tidy data. *Journal of Statistical Software*, 59(10):1–23, 2014. doi: 0.18637/jss.v059.i10. [p10]
- H. Wickham and R. Francois. *dplyr: A Grammar of Data Manipulation*, 2016. URL <https://CRAN.R-project.org/package=dplyr>. R package version 0.5.0. [p11]
- X.-G. Xin, T.-W. Wu, and Z. Jie. Introduction of CMIP5 experiments carried out with the climate system models of Beijing Climate Center. *Advances in Climate Change Research*, 4(1):41–49, 2013. doi: 10.3724/SPJ.1248.2013.041. [p8]

G. Brooke Anderson
Colorado State University
Department of Environmental & Radiological Health Sciences
1681 Campus Delivery
Fort Collins, Colorado 80523
brooke.anderson@colostate.edu

Colin Eason
Colorado State University
Department of Computer Science
1873 Campus Delivery
Fort Collins, Colorado 80523
aimesce@gmail.com

Elizabeth A. Barnes
Colorado State University
Department of Atmospheric Science
1371 Campus Delivery
Fort Collins, CO 80523
eabarnes@atmos.colostate.edu